

Quantitative Content Analysis: Lecture 2

Olga Kononykhina

15 Februar 2017

Today's Outline

Intro to R (Part I)























- R and RStudio
- Basic syntax
- Assigning objects
- Scalars, vectors, matrices
- Control flow: Loops and conditions

What is R?

R is an open source programming language with a particular focus on statistical programming

- Originally developed as 'S' by Bell Labs in 1993

R in comparison

Language Rank	Types	Spectrum Ranking
1. C	  	100.0
2. Java	  	98.1
3. Python	 	98.0
4. C++	  	95.9
5. R		87.9
6. C#	  	86.7
7. PHP		82.8
8. JavaScript	 	82.2
9. Ruby	 	74.5
10. Go	 	71.9

(source: IEEE Spectrum)

R interface (console-only)

- Command line, unlike SPSS/Stata
- An interpreted programming language
- Purist's interface: Text-editor & copy paste

RStudio is an Integrated Developer Environment (IDE) and serves as:

- Code editor
 - Code highlighting/completion, indentation, . . .
 - Feed code from editor to R-console
- Project manager
- Workspace viewer
- Data browser
- Enhanced output viewer
- Help browser

RStudio windows at startup

- Source editor
- Console
- Workspace
- Multi-purpose-panel

Projects

- A working directory for each project
- Code: .r files
- Dataset/Workspace: .Rdata files

Using RStudio (II)

Basic Workflow

- Edit in code editor (.r-file)
- Paste to console
- Save Workspace/Datasets (.Rdata-file)
- Save code routinely (no auto-save!)

Shortcuts

- Run code from editor: Select line and ctrl+Enter
- Switch between source and console: ctrl+1, ctrl+2
- Clear console: ctrl+L
- 'Arrow up' gives you the last line of code in the console
- Press Alt+Shift+K to see all keyboard shortcuts

Fundamentals of the R language

- Use # to comment code (will not be run)
- Case-sensitivity: data vs Data
- Assigning objects: <- and =

```
# Assign the number 5 to an object called number
```

```
number <- 5
```

```
number
```

```
## [1] 5
```

```
# Assign the character string Hello World
```

```
string <- "Hello World"
```

```
string
```

```
## [1] "Hello World"
```

Functions

Functions perform operations on the input given and end in ()

```
# e.g. find the class of number  
class(number)
```

```
## [1] "numeric"
```

Operations on scalars

You can use R as a calculator:

```
2 + 3  
2 - 3  
2 * 3  
2 / 3
```

Functions on scalars:

```
a <- 5  
factorial(a)
```

```
## [1] 120
```

Special values in R

- NA: not available, missing
- NULL: does not exist, is undefined
- TRUE, T: logical true
- FALSE, F: logical false

Finding special values

Function	Meaning
<code>is.na</code>	Is the value NA
<code>is.null</code>	Is the value NULL
<code>isTRUE</code>	Is the value TRUE
<code>!isTRUE</code>	Is the value FALSE

```
absent <- NA  
is.na(absent)
```

```
## [1] TRUE
```

Operations

Operator	Meaning
<	less than
>	greater than
==	equal to
<=	less than or equal to
>=	greater than or equal to
!=	not equal to
a b	a or b
a & b	a and b

Naming objects

- Object names cannot have spaces
 - Use CamelCase, `name_underscore`, or `name.period`
- Avoid creating an object with the same name as a function (e.g. `c` and `t`) or special value (`NA`, `NULL`, `TRUE`, `FALSE`).
- Use descriptive object names
- Each object name must be unique in an environment.
 - Assigning something to an object name that is already in use will overwrite the object's previous contents.

R is object-oriented

Objects are R's nouns and include (not exhaustive):

- character strings
- numbers
- vectors of numbers or character strings
- matrices
- data frames
- lists

Vectors

A vector is a container of objects put together in an order.

```
# Define a vector
```

```
a <- c(1,4,5)
```

```
b <- c(3,6,7)
```

```
# Join multiple vectors
```

```
ab <- c(a,b)
```

```
# Find vector length (number of its elements)
```

```
length(a)
```

```
# Reference vector components
```

```
ab[4] # Index one element in vector
```

```
ab[4:6] # Index several elements in a vector
```

```
ab[ab==6] # Index with condition
```

Operations on vectors

Operation	Meaning
<code>sort(x)</code>	sort a vector
<code>sum(x)</code>	sum of vector elements
<code>mean(x)</code>	arithmetic mean
<code>median(x)</code>	median value
<code>var(x)</code>	variance
<code>sd(x)</code>	standard deviation
<code>factorial(x)</code>	factorial of a number

Task 1

Calculate the mean of the vector 1,99,3,4,5,6,8. What's the mean if you out the 'outlier'?

Task 1 (solution)

Calculate the mean of the vector 1,99,3,4,5,6,8. What's the mean if you out the 'outlier'?

```
# Defining vector using sequence between 3 and 6
```

```
a <- c(1,99,3:6,8)
```

```
mean(a)
```

```
## [1] 18
```

```
# Calculate the mean of a but exclude the highest number
```

```
mean(a[a!=max(a)])
```

```
## [1] 4.5
```

Matrices

A Matrix is a square 2 dimensional container, i.e. vectors combined by row or column

- Must specify number of rows and columns

`matrix(x,nrow,ncol,byrow)`

- `x`: vector of length `nrow*ncol`
- `nrow`: number of rows
- `ncol`: number of columns
- `byrow`: TRUE or FALSE, specifies direction of input

Task 2

Assign a 6×10 matrix with $1, 2, 3, \dots, 60$ as the data.

Task 2 (solution)

Assign a 6×10 matrix with 1,2,3,...,60 as the data.

```
m <- matrix(1:60, nrow=6, ncol=10)
m
```

##		[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]
##	[1,]	1	7	13	19	25	31	37	43	49	55
##	[2,]	2	8	14	20	26	32	38	44	50	56
##	[3,]	3	9	15	21	27	33	39	45	51	57
##	[4,]	4	10	16	22	28	34	40	46	52	58
##	[5,]	5	11	17	23	29	35	41	47	53	59
##	[6,]	6	12	18	24	30	36	42	48	54	60

Referencing matrices

- Like vectors, you can reference matrices by elements
 - $a[i, j]$ refers to the i th row, j th column element of object a
- Can also reference rows/columns, these are vectors
 - $a[i,]$ is i th row, $a[, j]$ is j th column

Task 3

Extract the 9th column of the matrix from the previous problem. How can you find the 4th element in the 9th column?

Task 3 (solution)

Extract the 9th column of the matrix from the previous problem. How can you find the 4th element in the 9th column?

```
m[,9]
```

```
## [1] 49 50 51 52 53 54
```

```
m[4,9]
```

```
## [1] 52
```

```
m[,9][4]
```

```
## [1] 52
```

For() loops

For() loops are used to loop around a vector/matrix to do something.

```
m <- matrix(1:5, nrow=1, ncol=5)
m
```

```
##      [,1] [,2] [,3] [,4] [,5]
## [1,]    1    2    3    4    5
```

```
for (j in 1:3){
  m[,j]=0
}
m
```

```
##      [,1] [,2] [,3] [,4] [,5]
## [1,]    0    0    0    4    5
```

For() loops (II)

You can also 'nest' a for() loop in another for() loop

```
m <- matrix(1:15, nrow=3, ncol=5)
for (i in 1:2){
  for (j in 1:4){
    m[i,j]=0
  }
}
m
```

```
##      [,1] [,2] [,3] [,4] [,5]
## [1,]    0    0    0    0   13
## [2,]    0    0    0    0   14
## [3,]    3    6    9   12   15
```

Task 4

Create a matrix with `matrix(1:20,nrow=4,ncol=5)` and another with `matrix(NA,nrow=4,ncol=5)`. Adapt the second to the first matrix using `for()`

- hint: define a 'counter' variable that increments by 1 each time you move to the next cell

Task 4 (solution)

Create a matrix with `matrix(1:20,nrow=4,ncol=5)` and another with `matrix(NA,nrow=4,ncol=5)`. Adapt the second to the first matrix using `for()`

```
counter=1
m1 <- matrix(1:20,nrow=4,ncol=5)
m2 <- matrix(NA,nrow=4,ncol=5)
for (j in 1:5){
  for (i in 1:4){
    m2[i,j]=counter
    counter=counter+1
  }
}
```

Task 4 (alternative solution)

Create a matrix with `matrix(1:20,nrow=4,ncol=5)` and another with `matrix(NA,nrow=4,ncol=5)`. Adapt the second to the first matrix using `for()`

```
m1 <- matrix(1:20,nrow=4,ncol=5)
m2 <- matrix(NA,nrow=4,ncol=5)
for (j in 1:5){
  for (i in 1:4){
    m2[i,j]=m1[i,j]
  }
}
```

If() statements

If() statements are used to make conditions on executing some code. If condition is true, action is done.

```
a <- 3
b <- 4
number <- 0
if(a<b){
  number=number+1
}
number
```

```
## [1] 1
```

Tests for conditions: ==; >; <; >; <; !=

Task 6

Create the two objects `a <- sample(c(4,0),20,replace=TRUE)` and `m <- matrix(a,nrow=4,ncol=5)`. Now recode all the 4s into 1s using `if()` and `for()` statements.

Task 6 (solution)

Create the objects `a <- sample(c(4,0),20,replace=TRUE)` and `b <- matrix(a,nrow=4,ncol=5)`. Now recode all the 4s into 1s in `b` using `if()` and `for()` statements.

```
a <- sample(c(4,0),20,replace=TRUE)
b <- matrix(a,nrow=4,ncol=5)

for (j in 1:5){
  for (i in 1:4){
    if (b[i,j]==4){
      b[i,j]=1
    }
  }
}
```

- Type `help()` (or `?`) to see documentation
- Read Wickham & Grolemund's *R For Data Science Handbook*
- Check out the help function for a couple of functions used in today's course to see 'what is what' in the documentation

Next Session

- More sophisticated objects (dataframes, lists etc.)
- Basic plotting
- Reading and saving data
- Basic text-manipulation with R